

# car-price-prediction-1

December 22, 2023

```
[2]: # This Python 3 environment comes with many helpful analytics libraries
      ↪ installed
      # It is defined by the kaggle/python Docker image: https://github.com/kaggle/
      ↪ docker-python
      # For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt # data processing, CSV file I/O (e.g. pd.
      ↪ read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list
      ↪ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that
      ↪ gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved
      ↪ outside of the current session
```

```
[3]: df = pd.read_csv("F:/499a/CarPrice_Assignment.csv")
```

```
[4]: # Set display options to show all columns
      pd.set_option('display.max_columns', None)

      # Display the DataFrame
      df
```

```
[4]:
```

	car_ID	symboling	CarName	fueltype	aspiration	\
0	1	3	alfa-romero giulia	gas	std	
1	2	3	alfa-romero stelvio	gas	std	

2	3	1	alfa-romero	Quadrifoglio	gas	std
3	4	2		audi 100 ls	gas	std
4	5	2		audi 100ls	gas	std
..	...	...		...	...	...
200	201	-1		volvo 145e (sw)	gas	std
201	202	-1		volvo 144ea	gas	turbo
202	203	-1		volvo 244dl	gas	std
203	204	-1		volvo 246	diesel	turbo
204	205	-1		volvo 264gl	gas	turbo

	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	\
0	two	convertible	rwd	front	88.6	168.8	
1	two	convertible	rwd	front	88.6	168.8	
2	two	hatchback	rwd	front	94.5	171.2	
3	four	sedan	fwd	front	99.8	176.6	
4	four	sedan	4wd	front	99.4	176.6	
..	...	...	...	...	...	...	
200	four	sedan	rwd	front	109.1	188.8	
201	four	sedan	rwd	front	109.1	188.8	
202	four	sedan	rwd	front	109.1	188.8	
203	four	sedan	rwd	front	109.1	188.8	
204	four	sedan	rwd	front	109.1	188.8	

	carwidth	carheight	curbweight	enginetype	cylindernumber	enginesize	\
0	64.1	48.8	2548	dohc	four	130	
1	64.1	48.8	2548	dohc	four	130	
2	65.5	52.4	2823	ohcv	six	152	
3	66.2	54.3	2337	ohc	four	109	
4	66.4	54.3	2824	ohc	five	136	
..	...	...	...	...	...	...	
200	68.9	55.5	2952	ohc	four	141	
201	68.8	55.5	3049	ohc	four	141	
202	68.9	55.5	3012	ohcv	six	173	
203	68.9	55.5	3217	ohc	six	145	
204	68.9	55.5	3062	ohc	four	141	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	\
0	mpfi	3.47	2.68	9.0	111	5000	
1	mpfi	3.47	2.68	9.0	111	5000	
2	mpfi	2.68	3.47	9.0	154	5000	
3	mpfi	3.19	3.40	10.0	102	5500	
4	mpfi	3.19	3.40	8.0	115	5500	
..	...	...	...	...	...	...	
200	mpfi	3.78	3.15	9.5	114	5400	
201	mpfi	3.78	3.15	8.7	160	5300	
202	mpfi	3.58	2.87	8.8	134	5500	
203	idi	3.01	3.40	23.0	106	4800	

```
204      mpfi      3.78      3.15      9.5      114      5400
```

```
      citympg  highwaympg    price
0         21         27  13495.0
1         21         27  16500.0
2         19         26  16500.0
3         24         30  13950.0
4         18         22  17450.0
..      ...      ...      ...
200        23         28  16845.0
201        19         25  19045.0
202        18         23  21485.0
203        26         27  22470.0
204        19         25  22625.0
```

```
[205 rows x 26 columns]
```

```
[5]: df.columns
```

```
[5]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
          'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
          'carlength', 'carwidth', 'carheight', 'curbweight', 'engine-type',
          'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
          'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
          'price'],
          dtype='object')
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling             205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration            205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel           205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
```

```

13  curbweight      205 non-null    int64
14  enginetype      205 non-null    object
15  cylindernumber  205 non-null    object
16  enginesize      205 non-null    int64
17  fuelsystem      205 non-null    object
18  boreratio       205 non-null    float64
19  stroke          205 non-null    float64
20  compressionratio 205 non-null    float64
21  horsepower      205 non-null    int64
22  peakrpm         205 non-null    int64
23  citympg         205 non-null    int64
24  highwaympg      205 non-null    int64
25  price           205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
[7]: df.isnull().sum()
```

```

[7]: car_ID          0
     symboling      0
     CarName        0
     fueltype       0
     aspiration     0
     doornumber     0
     carbody        0
     drivewheel     0
     enginelocation 0
     wheelbase      0
     carlength      0
     carwidth       0
     carheight      0
     curbweight     0
     enginetype     0
     cylindernumber 0
     enginesize     0
     fuelsystem     0
     boreratio      0
     stroke         0
     compressionratio 0
     horsepower     0
     peakrpm        0
     citympg        0
     highwaympg     0
     price          0
     dtype: int64

```


```
[8]: df.duplicated().sum()
```

```
[8]: 0
```

```
[9]: plt.figure(figsize=(13, 9)) # Adjust the figure size as needed
sns.countplot(data=df, x="CarName", order=df["CarName"].value_counts().index)
plt.xticks(rotation=90, fontsize=10) # Rotate and adjust x-axis labels


# Adding labels and title
plt.xlabel("Car Name")
plt.ylabel("Frequency")
plt.title("Car Name Distribution")

# Show the plot
plt.tight_layout()
plt.show()
```



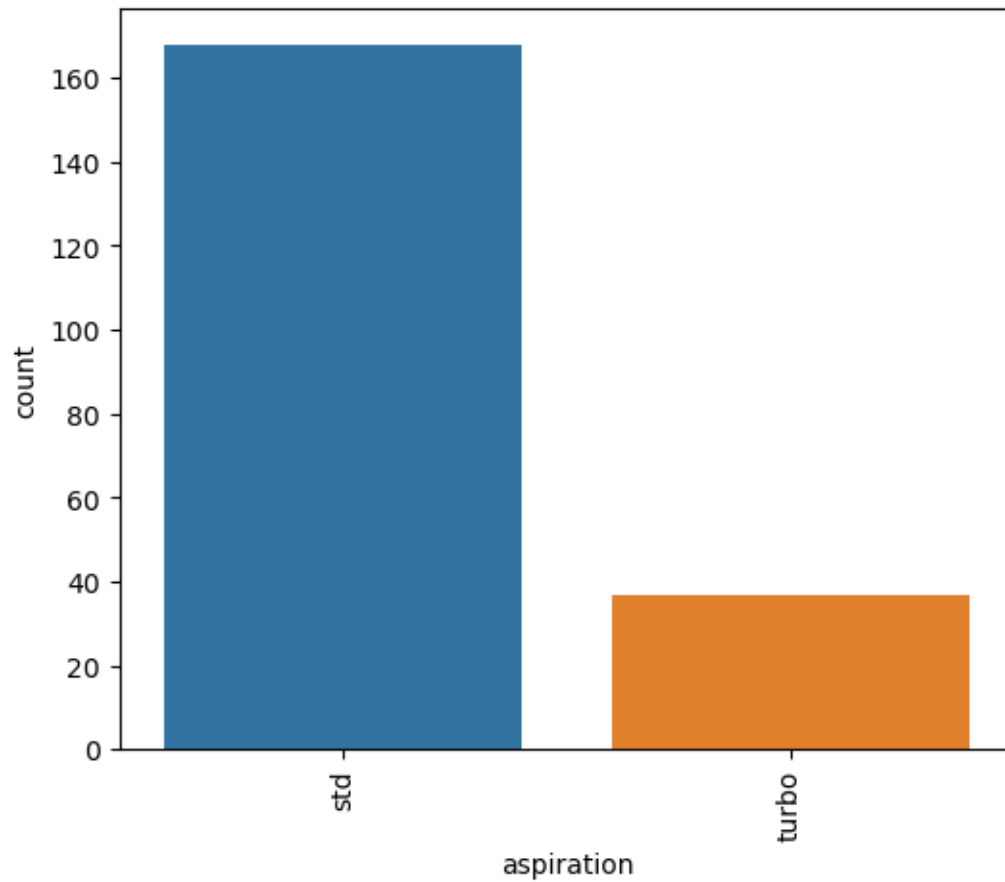
```
[10]: plt.figure(figsize=(6, 5)) # Adjust the figure size as needed
sns.countplot(data=df, x="fueltype", order=df["fueltype"].value_counts().index)
plt.xticks(rotation=90, fontsize=10) # Rotate and adjust x-axis labels
```

```
[10]: (array([0, 1]), [Text(0, 0, 'gas'), Text(1, 0, 'diesel')])
```




```
[11]: plt.figure(figsize=(6, 5)) # Adjust the figure size as needed
      sns.countplot(data=df, x="aspiration", order=df["aspiration"].value_counts().
      ↪index)
      plt.xticks(rotation=90, fontsize=10) # Rotate and adjust x-axis labels
```

```
[11]: (array([0, 1]), [Text(0, 0, 'std'), Text(1, 0, 'turbo')])
```



```
[12]: plt.figure(figsize=(6, 5)) # Adjust the figure size as needed
sns.countplot(data=df, x="doornumber", order=df["doornumber"].value_counts().
↪index)
plt.xticks(rotation=90, fontsize=10) # Rotate and adjust x-axis labels
```


```
[12]: (array([0, 1]), [Text(0, 0, 'four'), Text(1, 0, 'two')])
```



```
[13]: plt.figure(figsize=(6, 5)) # Adjust the figure size as needed
sns.countplot(data=df, x="carbody", order=df["carbody"].value_counts().index)
plt.xticks(rotation=90, fontsize=10) # Rotate and adjust x-axis labels
```


```
[13]: (array([0, 1, 2, 3, 4]),
      [Text(0, 0, 'sedan'),
        Text(1, 0, 'hatchback'),
        Text(2, 0, 'wagon'),
        Text(3, 0, 'hardtop'),
        Text(4, 0, 'convertible')])
```






```
[14]: plt.figure(figsize=(6, 5)) # Adjust the figure size as needed
sns.countplot(data=df, x="drivewheel", order=df["drivewheel"].value_counts().
    ↪index)
plt.xticks(rotation=90, fontsize=10) # Rotate and adjust x-axis labels
```

```
[14]: (array([0, 1, 2]), [Text(0, 0, 'fwd'), Text(1, 0, 'rwd'), Text(2, 0, '4wd')])
```



```
[15]: correlation_matrix = df.corr()

# Create a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Heatmap of Correlation Matrix')
plt.show()
```



```
[16]: df.head()
```

```
[16]:
```

	car_ID	symboling	CarName	fueltype	aspiration	door	number
0	1	3	alfa-romero giulia	gas	std	two	
1	2	3	alfa-romero stelvio	gas	std	two	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	
3	4	2	audi 100 ls	gas	std	four	
4	5	2	audi 100ls	gas	std	four	

	carbody	drivewheel	engine	location	wheelbase	carlength	carwidth
0	convertible	rwd	front		88.6	168.8	64.1
1	convertible	rwd	front		88.6	168.8	64.1
2	hatchback	rwd	front		94.5	171.2	65.5
3	sedan	fwd	front		99.8	176.6	66.2
4	sedan	4wd	front		99.4	176.6	66.4

	carheight	curbweight	enginetype	cylindernumber	enginesize	fuelsystem	\
0	48.8	2548	dohc	four	130	mpfi	
1	48.8	2548	dohc	four	130	mpfi	
2	52.4	2823	ohcv	six	152	mpfi	
3	54.3	2337	ohc	four	109	mpfi	
4	54.3	2824	ohc	five	136	mpfi	

	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
0	3.47	2.68	9.0	111	5000	21	
1	3.47	2.68	9.0	111	5000	21	
2	2.68	3.47	9.0	154	5000	19	
3	3.19	3.40	10.0	102	5500	24	
4	3.19	3.40	8.0	115	5500	18	












	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0

```
[17]: def scatter(x, title, fig):
    plt.subplot(6, 2, fig)
    plt.scatter(df[x], df['price'])
    plt.title(title + ' vs Price')
    plt.ylabel('Price')
    plt.xlabel(x)

plt.figure(figsize=(10, 20))

scatter('wheelbase', 'Wheelbase', 1)
scatter('carlength', 'Car Length', 2)
scatter('carheight', 'Car Height', 3)
scatter('curbweight', 'Curb Weight', 4)
scatter('enginesize', 'Engine Size', 5)
scatter('boreratio', 'Bore Ratio', 6)
scatter('stroke', 'Stroke', 7)
scatter('horsepower', 'Horsepower', 8)
scatter('peakrpm', 'Peak RPM', 9)
scatter('highwaympg', 'Highway MPG', 10)
scatter('citympg', 'City MPG', 11) # Change '11' to '1' or '2' based on your
    ↪ desired position

plt.tight_layout()
plt.show()
```



```
[18]: df.drop(columns=["car_ID", "peakrpm", "stroke", "carheight"], inplace=True)
```

```
[19]: from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
[20]: en = LabelEncoder()
catCols = [
    ↪ ['CarName', 'fueltype', 'aspiration', 'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'engine', 'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'compressionratio', 'horsepower', 'citympg', 'highwaympg', 'price']
for cols in catCols:
    df[cols] = en.fit_transform(df[cols])
```

```
[21]: df.head()
```

```
[21]:
```

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	\
0	3	2	1	0	1	0	2	
1	3	3	1	0	1	0	2	
2	1	1	1	0	1	2	2	
3	2	4	1	0	0	3	1	
4	2	5	1	0	0	3	0	

	enginelocation	wheelbase	carlength	carwidth	curbweight	enginetype	\
0	0	88.6	168.8	64.1	2548	0	
1	0	88.6	168.8	64.1	2548	0	
2	0	94.5	171.2	65.5	2823	5	
3	0	99.8	176.6	66.2	2337	3	
4	0	99.4	176.6	66.4	2824	3	

	cylindernumber	enginesize	fuelsystem	boreratio	compressionratio	\
0	2	130	5	3.47	9.0	
1	2	130	5	3.47	9.0	
2	3	152	5	2.68	9.0	
3	2	109	5	3.19	10.0	
4	1	136	5	3.19	8.0	

	horsepower	citympg	highwaympg	price
0	111	21	27	13495.0
1	111	21	27	16500.0
2	154	19	26	16500.0
3	102	24	30	13950.0
4	115	18	22	17450.0

```
[22]: X = df.drop("price", axis = 1)
y = df["price"]
```

```
[23]: X.head()
```

```
[23]:   symboling  CarName  fueltype  aspiration  doornumber  carbody  drivewheel  \
0         3         2         1           0           1         0           2
1         3         3         1           0           1         0           2
2         1         1         1           0           1         2           2
3         2         4         1           0           0         3           1
4         2         5         1           0           0         3           0

   enginelocation  wheelbase  carlength  carwidth  curbweight  enginetype  \
0                0      88.6      168.8      64.1       2548         0
1                0      88.6      168.8      64.1       2548         0
2                0      94.5      171.2      65.5       2823         5
3                0      99.8      176.6      66.2       2337         3
4                0      99.4      176.6      66.4       2824         3

   cylindernumber  enginesize  fuelsystem  boreratio  compressionratio  \
0                2         130          5        3.47             9.0
1                2         130          5        3.47             9.0
2                3         152          5        2.68             9.0
3                2         109          5        3.19            10.0
4                1         136          5        3.19             8.0

   horsepower  citympg  highwaympg
0          111       21          27
1          111       21          27
2          154       19          26
3          102       24          30
4          115       18          22
```

```
[24]: y.head()
```

```
[24]: 0    13495.0
1    16500.0
2    16500.0
3    13950.0
4    17450.0
Name: price, dtype: float64
```

```
[25]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.1,random_state=4)
from sklearn.preprocessing import StandardScaler

# Create a MinMaxScaler instance
scaler = StandardScaler()

# Fit the scaler on the training data and transform it
X_train = scaler.fit_transform(X_train)
```

```
# Transform the test data using the same scaler
X_test = scaler.transform(X_test)
```

```
[32]: from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import Ridge
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor,
↳ BaggingRegressor, ExtraTreesRegressor, GradientBoostingRegressor
from xgboost import XGBRegressor
from sklearn.metrics import accuracy_score
```

```
[36]: def train_regressor(regressor, X_train, y_train, X_test, y_test):
    regressor.fit(X_train, y_train)
    y_pred = regressor.predict(X_test)

    r2 = regressor.score(X_test, y_test)
    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)

    return r2, mse, mae
regressors = {
    "SVR": SVR(kernel='linear', C=1),
    "KNeighborsRegressor": KNeighborsRegressor(n_neighbors=5),
    "DecisionTreeRegressor": DecisionTreeRegressor(max_depth=5),
    "Ridge": Ridge(alpha=1.0),
    "RandomForestRegressor": RandomForestRegressor(n_estimators=200, max_depth=3,
↳ random_state=2),
    "AdaBoostRegressor": AdaBoostRegressor(n_estimators=100, random_state=2),
    "BaggingRegressor": BaggingRegressor(n_estimators=100, random_state=2),
    "ExtraTreesRegressor": ExtraTreesRegressor(n_estimators=100,
↳ random_state=2),
    "GradientBoostingRegressor": GradientBoostingRegressor(learning_rate= 0.2,
↳ min_samples_leaf= 4, min_samples_split= 5, n_estimators=100),
    "XGBRegressor": XGBRegressor(n_estimators=100, random_state=2)
}
```

```
[48]: from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import Ridge
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor,
↳ BaggingRegressor, ExtraTreesRegressor, GradientBoostingRegressor
```



```

from xgboost import XGBRegressor
import time

def train_regressor(regressor, X_train, y_train, X_test, y_test):
    training_times = [] # List to store training times for each epoch
    for epoch in range(epochs):
        start_time = time.time() # Record the start time of the epoch
        regressor.fit(X_train, y_train)
        end_time = time.time() # Record the end time of the epoch
        epoch_time = end_time - start_time # Calculate epoch training time
        training_times.append(epoch_time)

        y_pred = regressor.predict(X_test)
        r2 = regressor.score(X_test, y_test)
        mse = mean_squared_error(y_test, y_pred)
        mae = mean_absolute_error(y_test, y_pred)

        print(f"Epoch {epoch + 1} - R-squared: {r2:.2f}, MSE: {mse:.2f}, MAE: {mae:.2f}, Time: {epoch_time:.2f} seconds")

    return training_times

regressors = {
    "SVR": SVR(kernel='linear', C=1),
    "KNeighborsRegressor": KNeighborsRegressor(n_neighbors=5),
    "DecisionTreeRegressor": DecisionTreeRegressor(max_depth=5),
    "Ridge": Ridge(alpha=1.0),
    "RandomForestRegressor": RandomForestRegressor(n_estimators=200,
    ↪max_depth=3, random_state=2),
    "AdaBoostRegressor": AdaBoostRegressor(n_estimators=100, random_state=2),
    "BaggingRegressor": BaggingRegressor(n_estimators=100, random_state=2),
    "ExtraTreesRegressor": ExtraTreesRegressor(n_estimators=100,
    ↪random_state=2),
    "GradientBoostingRegressor": GradientBoostingRegressor(learning_rate=0.2,
    ↪min_samples_leaf=4, min_samples_split=5, n_estimators=100),
    "XGBRegressor": XGBRegressor(n_estimators=100, random_state=2)
}

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=4) # Split your dataset into training and testing sets
epochs = 10 # Set the number of epochs

results = {} # To store training times for each regressor

for regressor_name, regressor in regressors.items():
    print(f"Training {regressor_name}...")

```

```

    training_times = train_regressor(regressor, X_train, y_train, X_test,
    ↪y_test)
    results[regressor_name] = training_times

# Print the training times for each epoch for each regressor
for regressor_name, training_times in results.items():
    print(f"{regressor_name} Training Times (in seconds): {training_times}")

```

Training SVR...

```

Epoch 1 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.02 seconds
Epoch 2 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 3 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 4 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 5 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 6 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 7 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 8 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 9 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds
Epoch 10 - R-squared: 0.81, MSE: 9494398.64, MAE: 2307.18, Time: 0.01 seconds

```

Training KNeighborsRegressor...

```

Epoch 1 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 2 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 3 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 4 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 5 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 6 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 7 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 8 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 9 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds
Epoch 10 - R-squared: 0.74, MSE: 12723427.90, MAE: 2361.76, Time: 0.00 seconds

```

Training DecisionTreeRegressor...

```

Epoch 1 - R-squared: 0.88, MSE: 6163271.96, MAE: 1892.47, Time: 0.00 seconds
Epoch 2 - R-squared: 0.88, MSE: 5995462.40, MAE: 1830.23, Time: 0.00 seconds
Epoch 3 - R-squared: 0.87, MSE: 6577179.04, MAE: 1941.35, Time: 0.00 seconds
Epoch 4 - R-squared: 0.85, MSE: 7309910.35, MAE: 2018.61, Time: 0.00 seconds
Epoch 5 - R-squared: 0.88, MSE: 6163271.96, MAE: 1892.47, Time: 0.00 seconds
Epoch 6 - R-squared: 0.87, MSE: 6455411.05, MAE: 1941.35, Time: 0.00 seconds
Epoch 7 - R-squared: 0.85, MSE: 7278440.34, MAE: 2002.93, Time: 0.00 seconds
Epoch 8 - R-squared: 0.85, MSE: 7263868.78, MAE: 1956.37, Time: 0.00 seconds
Epoch 9 - R-squared: 0.88, MSE: 6163271.96, MAE: 1892.47, Time: 0.00 seconds
Epoch 10 - R-squared: 0.87, MSE: 6455411.05, MAE: 1941.35, Time: 0.00 seconds

```

Training Ridge...

```

Epoch 1 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds
Epoch 2 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds
Epoch 3 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds
Epoch 4 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds
Epoch 5 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds

```

Epoch 6 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds  
Epoch 7 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds  
Epoch 8 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds  
Epoch 9 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds  
Epoch 10 - R-squared: 0.83, MSE: 8196426.81, MAE: 2170.84, Time: 0.00 seconds  
Training RandomForestRegressor...

Epoch 1 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.36 seconds  
Epoch 2 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.30 seconds  
Epoch 3 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.30 seconds  
Epoch 4 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.28 seconds  
Epoch 5 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.28 seconds  
Epoch 6 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.28 seconds  
Epoch 7 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.29 seconds  
Epoch 8 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.28 seconds  
Epoch 9 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.29 seconds  
Epoch 10 - R-squared: 0.87, MSE: 6582210.91, MAE: 1914.52, Time: 0.28 seconds  
Training AdaBoostRegressor...

Epoch 1 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.18 seconds  
Epoch 2 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.15 seconds  
Epoch 3 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.17 seconds  
Epoch 4 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.16 seconds  
Epoch 5 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.16 seconds  
Epoch 6 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.18 seconds  
Epoch 7 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.16 seconds  
Epoch 8 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.14 seconds  
Epoch 9 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.19 seconds  
Epoch 10 - R-squared: 0.89, MSE: 5259872.25, MAE: 1933.86, Time: 0.15 seconds  
Training BaggingRegressor...

Epoch 1 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.38 seconds  
Epoch 2 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.43 seconds  
Epoch 3 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.53 seconds  
Epoch 4 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.47 seconds  
Epoch 5 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.39 seconds  
Epoch 6 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.37 seconds  
Epoch 7 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.37 seconds  
Epoch 8 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.35 seconds  
Epoch 9 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.35 seconds  
Epoch 10 - R-squared: 0.91, MSE: 4226915.74, MAE: 1486.30, Time: 0.36 seconds  
Training ExtraTreesRegressor...

Epoch 1 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.21 seconds  
Epoch 2 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.23 seconds  
Epoch 3 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.22 seconds  
Epoch 4 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.20 seconds  
Epoch 5 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.21 seconds  
Epoch 6 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.21 seconds  
Epoch 7 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.22 seconds  
Epoch 8 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.21 seconds  
Epoch 9 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.21 seconds

Epoch 10 - R-squared: 0.92, MSE: 4074228.75, MAE: 1627.64, Time: 0.23 seconds  
 Training GradientBoostingRegressor...

Epoch 1 - R-squared: 0.93, MSE: 3420535.61, MAE: 1464.40, Time: 0.12 seconds  
 Epoch 2 - R-squared: 0.93, MSE: 3410006.06, MAE: 1463.99, Time: 0.11 seconds  
 Epoch 3 - R-squared: 0.93, MSE: 3501758.93, MAE: 1482.52, Time: 0.13 seconds  
 Epoch 4 - R-squared: 0.93, MSE: 3415039.28, MAE: 1453.30, Time: 0.12 seconds  
 Epoch 5 - R-squared: 0.93, MSE: 3424994.32, MAE: 1465.81, Time: 0.13 seconds  
 Epoch 6 - R-squared: 0.93, MSE: 3406715.27, MAE: 1459.56, Time: 0.15 seconds  
 Epoch 7 - R-squared: 0.93, MSE: 3542737.97, MAE: 1495.00, Time: 0.12 seconds  
 Epoch 8 - R-squared: 0.93, MSE: 3427623.50, MAE: 1466.08, Time: 0.15 seconds  
 Epoch 9 - R-squared: 0.93, MSE: 3419481.91, MAE: 1454.70, Time: 0.12 seconds  
 Epoch 10 - R-squared: 0.93, MSE: 3558609.02, MAE: 1489.47, Time: 0.12 seconds  
 Training XGBRegressor...

Epoch 1 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.10 seconds  
 Epoch 2 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.09 seconds  
 Epoch 3 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.12 seconds  
 Epoch 4 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.09 seconds  
 Epoch 5 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.09 seconds  
 Epoch 6 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.08 seconds  
 Epoch 7 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.09 seconds  
 Epoch 8 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.09 seconds  
 Epoch 9 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.09 seconds  
 Epoch 10 - R-squared: 0.92, MSE: 3711713.11, MAE: 1324.63, Time: 0.09 seconds  
 SVR Training Times (in seconds): [0.015958786010742188, 0.013962268829345703, 0.012964725494384766, 0.013965368270874023, 0.0139617919921875, 0.013960838317871094, 0.01302337646484375, 0.012023448944091797, 0.012023448944091797, 0.013007402420043945]  
 KNeighborsRegressor Training Times (in seconds): [0.0020248889923095703, 0.0022466182708740234, 0.0009970664978027344, 0.0009970664978027344, 0.001994609832763672, 0.0009970664978027344, 0.000997304916381836, 0.0009968280792236328, 0.0009965896606445312, 0.0019936561584472656]  
 DecisionTreeRegressor Training Times (in seconds): [0.0019941329956054688, 0.001994609832763672, 0.0029840469360351562, 0.0019948482513427734, 0.001994609832763672, 0.001993894577026367, 0.001994609832763672, 0.001994609832763672, 0.002992868423461914, 0.002991914749145508]  
 Ridge Training Times (in seconds): [0.0019958019256591797, 0.0019943714141845703, 0.001995086669921875, 0.001994609832763672, 0.0019953250885009766, 0.001994609832763672, 0.0019943714141845703, 0.002006053924560547, 0.0019943714141845703, 0.0009975433349609375]  
 RandomForestRegressor Training Times (in seconds): [0.35610437393188477, 0.2971975803375244, 0.29517292976379395, 0.282275915145874, 0.2792532444000244, 0.28029561042785645, 0.28623199462890625, 0.28224754333496094, 0.29026293754577637, 0.27829551696777344]  
 AdaBoostRegressor Training Times (in seconds): [0.18350958824157715, 0.15259027481079102, 0.16655611991882324, 0.16161036491394043, 0.1585242748260498, 0.17547845840454102, 0.15953278541564941, 0.14162087440490723, 0.18550562858581543, 0.15255141258239746]  
 BaggingRegressor Training Times (in seconds): [0.3809685707092285,

```

0.43088865280151367, 0.5305945873260498, 0.474729061126709, 0.38999199867248535,
0.3660435676574707, 0.3700551986694336, 0.35400986671447754,
0.35301899909973145, 0.36305880546569824]
ExtraTreesRegressor Training Times (in seconds): [0.20939874649047852,
0.23342037200927734, 0.21642446517944336, 0.20345544815063477,
0.20647621154785156, 0.20644760131835938, 0.22042274475097656,
0.20648550987243652, 0.21342992782592773, 0.22838807106018066]
GradientBoostingRegressor Training Times (in seconds): [0.11667943000793457,
0.11170077323913574, 0.133683443069458, 0.12067842483520508,
0.12865281105041504, 0.14561057090759277, 0.12462282180786133,
0.14760732650756836, 0.12462782859802246, 0.12366008758544922]
XGBRegressor Training Times (in seconds): [0.0967404842376709,
0.08976054191589355, 0.11768627166748047, 0.09076642990112305,
0.09474587440490723, 0.0827779769897461, 0.09275102615356445,
0.09075665473937988, 0.08676767349243164, 0.09075760841369629]

```

```

[37]: # Lists to store metrics
algorithm_names = []
r2_scores = []
mse_scores = []
mae_scores = []

# Train and evaluate each regressor
for name, regressor in regressors.items():
    r2, mse, mae= train_regressor(regressor, X_train, y_train, X_test, y_test)

    algorithm_names.append(name)
    r2_scores.append(r2)
    mse_scores.append(mse)
    mae_scores.append(mae)

# Create a DataFrame to store the results
results_df = pd.DataFrame({
    'Algorithm': algorithm_names,
    'R2 Score': r2_scores,
    'Mean Squared Error': mse_scores,
    'Mean Absolute Error': mae_scores
})

display(results_df)

```

	Algorithm	R2 Score	Mean Squared Error \
0	SVR	-0.121589	5.598468e+07
1	KNeighborsRegressor	0.857828	7.096604e+06
2	DecisionTreeRegressor	0.847318	7.621191e+06
3	Ridge	0.846768	7.648635e+06
4	RandomForestRegressor	0.833496	8.311119e+06

5	AdaBoostRegressor	0.899744	5.004332e+06
6	BaggingRegressor	0.904408	4.771540e+06
7	ExtraTreesRegressor	0.911473	4.418865e+06
8	GradientBoostingRegressor	0.941728	2.908671e+06
9	XGBRegressor	0.913827	4.301364e+06

	Mean Absolute Error
0	5266.610617
1	2063.990476
2	2020.238631
3	2220.268551
4	2027.620255
5	1684.996219
6	1305.606825
7	1565.243414
8	1262.223506
9	1173.013346

```
[39]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score, \
      ↪ accuracy_score
      from sklearn.svm import SVR
      from sklearn.neighbors import KNeighborsRegressor
      from sklearn.tree import DecisionTreeRegressor
      from sklearn.linear_model import Ridge
      from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor, \
      ↪ BaggingRegressor, ExtraTreesRegressor, GradientBoostingRegressor
      from xgboost import XGBRegressor
      import numpy as np

      # Define your training and testing data: X_train, y_train, X_test, y_test

      threshold = 0.5 # You can adjust the threshold as needed

      for model_name, regressor in regressors.items():
          regressor.fit(X_train, y_train)
          y_pred = regressor.predict(X_test)

          # Binarize the predictions based on the threshold
          y_pred_binary = (y_pred >= threshold).astype(int)

          # Calculate regression metrics
          r2 = r2_score(y_test, y_pred)
          mse = mean_squared_error(y_test, y_pred)
          mae = mean_absolute_error(y_test, y_pred)

          # Calculate accuracy based on binarized predictions
          accuracy = accuracy_score(y_test, y_pred_binary)
```

```
print(f"{model_name} - R-squared: {r2}, MSE: {mse}, MAE: {mae}, Accuracy: {accuracy}")
```

```
SVR - R-squared: -0.12158881504099561, MSE: 55984678.84228565, MAE:
5266.61061658763, Accuracy: 0.0
KNeighborsRegressor - R-squared: 0.8578276878993255, MSE: 7096603.60952381, MAE:
2063.9904761904763, Accuracy: 0.0
DecisionTreeRegressor - R-squared: 0.853439936266576, MSE: 7315620.474443436,
MAE: 1906.5522954796409, Accuracy: 0.0
Ridge - R-squared: 0.8467683775631772, MSE: 7648634.736335464, MAE:
2220.268551296449, Accuracy: 0.0
RandomForestRegressor - R-squared: 0.8334962668337302, MSE: 8311118.925534684,
MAE: 2027.6202545102465, Accuracy: 0.0
AdaBoostRegressor - R-squared: 0.8997439443063169, MSE: 5004332.251476664, MAE:
1684.9962185038569, Accuracy: 0.0
BaggingRegressor - R-squared: 0.9044076645171308, MSE: 4771540.273961137, MAE:
1305.606825238095, Accuracy: 0.0
ExtraTreesRegressor - R-squared: 0.911473104993635, MSE: 4418865.202088281, MAE:
1565.243413809524, Accuracy: 0.0
GradientBoostingRegressor - R-squared: 0.9382325041087549, MSE:
3083156.121022084, MAE: 1297.3170644570703, Accuracy: 0.0
XGBRegressor - R-squared: 0.9138271000295217, MSE: 4301364.336964516, MAE:
1173.0133463541667, Accuracy: 0.0
```

```
[41]: import pandas as pd

# Lists to store metrics
algorithm_names = []
r2_scores = []
mse_scores = []
mae_scores = []
accuracy_scores = [] # Updated this variable name to avoid conflicts

# Train and evaluate each regressor
for model_name, regressor in regressors.items():
    r2, mse, mae = train_regressor(regressor, X_train, y_train, X_test, y_test)

    algorithm_names.append(model_name) # Corrected variable name
    r2_scores.append(r2)
    mse_scores.append(mse)
    mae_scores.append(mae)
    accuracy_scores.append(accuracy) # You need to calculate accuracy here

results_df = pd.DataFrame({
    'Algorithm': algorithm_names,
    'R2 Score': r2_scores,
```

```

    'Mean Squared Error': mse_scores,
    'Mean Absolute Error': mae_scores,
    'Accuracy': accuracy_scores
})

# Display the DataFrame
display(results_df)

```

	Algorithm	R2 Score	Mean Squared Error \
0	SVR	-0.121589	5.598468e+07
1	KNeighborsRegressor	0.857828	7.096604e+06
2	DecisionTreeRegressor	0.847318	7.621191e+06
3	Ridge	0.846768	7.648635e+06
4	RandomForestRegressor	0.833496	8.311119e+06
5	AdaBoostRegressor	0.899744	5.004332e+06
6	BaggingRegressor	0.904408	4.771540e+06
7	ExtraTreesRegressor	0.911473	4.418865e+06
8	GradientBoostingRegressor	0.946082	2.691336e+06
9	XGBRegressor	0.913827	4.301364e+06

	Mean Absolute Error	Accuracy
0	5266.610617	[]
1	2063.990476	[]
2	2020.238631	[]
3	2220.268551	[]
4	2027.620255	[]
5	1684.996219	[]
6	1305.606825	[]
7	1565.243414	[]
8	1180.561497	[]
9	1173.013346	[]